

Performance-Estimation Properties of Cross-Validation-Based Protocols with Simultaneous Hyper-Parameter Optimization

Ioannis Tsamardinos^{1,2}, Amin Rakhshani^{1,2}, Vincenzo Lagani¹

¹Institute of Computer Science, Foundation for Research and Technology Hellas, Heraklion, Greece

²Computer Science Department, University of Crete, Heraklion, Greece

{ tsamard, vlagani, aminra }@ics.forth.gr

Abstract. In a typical supervised data analysis task, one needs to perform the following two tasks: (a) select the best combination of learning methods (e.g., for variable selection and classifier) and tune their hyper-parameters (e.g., K in K -NN), also called *model selection*, and (b) provide an estimate of the performance of the final, reported model. Combining the two tasks is not trivial because when one selects the set of hyper-parameters that seem to provide the best estimated performance, this estimation is optimistic (biased / overfitted) due to performing multiple statistical comparisons. In this paper, we confirm that the simple Cross-Validation with model selection is indeed optimistic (overestimates) in small sample scenarios. In comparison the Nested Cross Validation and the method by Tibshirani and Tibshirani provide conservative estimations, with the later protocol being more computationally efficient. The role of stratification of samples is examined and it is shown that stratification is beneficial.

1 Introduction

A typical supervised analysis (e.g., classification or regression) consists of several steps that result in a final, single prediction, or diagnostic model. For example, the analyst may need to impute missing values, perform variable selection or general dimensionality reduction, discretize variables, try several different representations of the data, and finally, apply a learning algorithm for classification or regression. Each of these steps requires a selection of algorithms out of hundreds or even thousands of possible choices, as well as the tuning of their hyper-parameters¹. *Hyper-parameter*

¹ We use the term “hyper-parameters” to denote the algorithm parameters that can be set by the user and are not estimated directly from the data, e.g., the parameter K in the K -NN algorithm. In contrast, the term “parameters” in the statistical literature typically refers to the

optimization is also called the *model selection* problem since each combination of hyper-parameters tried leads to a possible classification or regression model out of which the best is to be selected.

There are several alternatives in the literature about how to identify a good combination of methods and their hyper-parameters (e.g., [1][2]) and they all involve implicitly or explicitly searching the space of hyper-parameters and trying different combinations. Unfortunately, trying multiple combinations, estimating their performance, and *reporting the performance of the best model found leads to overestimating the performance (i.e., underestimate the error / loss)*, sometimes also referred to as overfitting². This phenomenon is called the problem of *multiple comparisons in induction algorithms* and has been analyzed in detail in [3] and is related to the *multiple testing* or *multiple comparisons* in statistical hypothesis testing. Intuitively, when one selects among several models whose estimations vary around their true mean value, it becomes likely that what seems to be the best model has been “lucky” in the specific test set and its performance has been overestimated. Extensive discussions and experiments on the subject can be found in [2].

The bias should increase with the number of models tried and decrease with the size of the test set. Notice that, when using Cross Validation-based protocols to estimate performance each sample serves once and only once as a test case. Thus, *one can consider the total data-set sample size as the size of the test set*. Typical high-dimensional datasets in biology often contain less than 100 samples and thus, one should be careful with the estimation protocols employed for their analysis.

What about the number of different models tried in an analysis? Is it realistic to expect an analyst to generate thousands of different models? Obviously, it is very rare that any analyst will employ thousands of different algorithms; however, most learning algorithms are parameterized by several different hyper-parameters. For example, the standard 1-norm, polynomial Support Vector Machine algorithm takes as hyper-parameters the cost C of misclassifications and the degree of the polynomial d . Similarly, most variable selection methods take as input a statistical significance threshold or the number of variables to return. If an analyst tries several different methods for imputation, discretization, variable selection, and classification, each with several different hyper-parameter values, the number of combinations explodes and can easily reach into the thousands.

model quantities that are estimated directly by the data, e.g., the weight vector w in a linear regression model $y = w \cdot x + b$. See [2] for a definition and discussion too.

² The term “overfitting” is a more general term and we prefer the term “overestimating” to characterize this phenomenon.

Algorithm 1: K-Fold Cross-Validation $\underline{CV}(D)$

Input: A dataset $D = \{\langle x_i, y_i \rangle, i=1, \dots, N\}$

Output: A model M

An estimation of performance (loss) of M

Randomly Partition D to K folds F_i

Model $M = f(\cdot, D)$ // the model learned on all data D

Estimation \widehat{L}_{CV} :

$$\widehat{e}_i = \frac{1}{N_i} \sum_{j \in F_i} L(y_j, f(x_j, D_{\setminus i})) , \quad \widehat{L}_{CV} = \frac{1}{K} \sum_{i=1}^K \widehat{e}_i$$

Return $\langle M, \widehat{L}_{CV} \rangle$

Notice that, model selection and optimistic estimation of performance *may also happen unintentionally and implicitly* in many other settings. More specifically, consider a typical publication where a new algorithm is introduced and its performance (after tuning the hyper-parameters) is compared against numerous other alternatives from the literature (again, after tuning their hyper-parameters), on several datasets. The comparison aims to comparatively evaluate the methods. However, *the reported performances of the best method on each dataset suffer from the same problem of multiple inductions and are on average optimistically estimated.*

In the remainder of the paper, we revisit the Cross-Validation (CV) protocol. We corroborate [2][4] that CV *overestimates* performance when it is used with hyper-parameter optimization. As expected overestimation of performance increases with decreasing sample sizes. We present two other performance estimation methods in the literature. The method by Tibshirani and Tibshirani (hereafter TT) [5] tries to estimate the bias and remove it from the estimation. The Nested Cross Validation (NCV) method [6] cross-validates the whole hyper-parameter optimization procedure (which includes an inner cross-validation, hence the name). NCV is a generalization of the technique where data is partitioned in train-validation-test sets. We show that both of them are conservative (underestimate) performance, while TT is computationally more efficient. To our knowledge, this is the first time the three methods are compared against each other on real datasets. *The excellent behavior of TT in these preliminary results makes it a promising alternative to NCV.*

The effect of stratification is also empirically examined. Stratification is a technique that during partitioning of the data into folds for cross-validation forces the same distribution of the outcome classes to each fold. When data are split randomly, on average, the distribution of the outcome in each fold will be the same as in the whole dataset. However, in small sample sizes or imbalanced data it could happen that a fold gets no samples that belong in one of the classes (or in general, the class distribution in a fold is very different from the original). Stratification ensures that this doesn't occur. We show that stratification decreases the variance of the estimation and thus should always be applied.

2 Cross-Validation Without Hyper-Parameter Optimization (CV)

K-fold Cross Validation is perhaps the most common method for estimating performance of a learning method for small and medium sample sizes. Despite its popularity, its theoretical properties are arguably not well known especially outside the machine learning community, particularly when it is employed with simultaneous hyper-parameter optimization, as evidenced by the following common machine learning books: Duda ([7], p. 484) presents CV without discussing it in the context of model selection and only hints that it may underestimate (when used without model selection): “The jackknife [i.e., leave-one-out CV] in particular, generally gives good estimates because each of the n classifiers is quite similar to the classifier being tested ...”. Similarly, Mitchell [8](pp. 112, 147, 150) mentions CV but in the context of hyper-parameter optimization. Bishop [9] does not deal at all with issues of performance estimation and model selection. A notable exception is the Hastie and co-authors [10] book that offers the best treatment of the subject, *upon which the following sections are based*. Yet, CV is still not discussed in the context of model selection.

Let’s assume a dataset $D = \{ \langle x_i, y_i \rangle, i=1, \dots, N \}$, of identically and independently distributed (i.i.d.) predictor vectors x_i and corresponding outcomes y_i . Let us also assume that we have a *single method* for learning from such data and producing a single prediction model. We will denote with $f(x_i, D)$ the output of the model produced by the learner f when trained on data D and applied on input x_i . The actual model produced by f on dataset D is denoted with $f(\cdot, D)$. We will denote with $L(y, y')$ the loss (error) measure of prediction y' when the true output is y . One common loss function is the zero-one loss function: $L(y, y') = 1$, if $y \neq y'$ and $L(y, y') = 0$, otherwise.

Thus, the average zero-one loss of a classifier equals 1-accuracy, i.e., it is the probability of making an incorrect classification. K-fold CV partitions the data D into K subsets called folds F_1, \dots, F_k . We denote with $D_{\setminus i}$ the data excluding fold F_i and N_i the sample size of each fold. The K-fold CV algorithm is shown in Algorithm 1.

First, notice that CV returns the model learned from all data D , $f(\cdot, D)$. This is the model to be employed operationally for classification. It then tries to estimate the performance of the returned model by constructing K other models from datasets $D_{\setminus i}$, each time excluding one fold from the training set. Each of these models is then applied on each fold F_i serving as test and the loss is averaged over all samples.

Is \widehat{L}_{CV} an unbiased estimate of the loss of $f(\cdot, D)$? First, notice that each sample x_i is used once and only once as a test case. Thus, effectively there are as many i.i.d. test cases as samples in the dataset. Perhaps, this characteristic is what makes the CV so popular versus other protocols such as repeatedly partitioning the dataset to train-test subsets. The test size being as large as possible could facilitate the estimation of the loss and its variance (although, theoretical results show that there is no unbiased estimator of the variance for the CV! [11]). However, test cases are predicted with different models! If these models were trained on independent train sets of size equal to the original data D , then CV would indeed estimate the average loss of the models pro-

duced by the specific learning method on the specific task when trained with the specific sample size. As it stands though, since the models are correlated and have smaller size than the original:

K-Fold CV estimates the average loss of models returned by the specific learning method f on the specific classification task when trained with subsets of D of size $|D_{\setminus i}|$

Since $|D_{\setminus i}| = (K-1)/K \cdot |D| < |D|$ (e.g., for 5-fold, we are using 80% of the total sample size for training each time) and *assuming that the learning method improves on average with larger sample sizes* we expect \widehat{L}_{CV} to be conservative (i.e., the true performance be underestimated). How conservative it will be depends on where the classifier is operating on its learning curve for this specific task. It also depends on the number of folds K : the larger the K , the more $(K-1)/K$ approaches 100% and the bias disappears, i.e., leave-one-out CV should be the least biased (however, there may be still be significant estimation problems, see [12], p. 151, and [4] for an extreme failure of leave-one-out CV). When sample sizes are small or distributions are imbalanced (i.e., some classes are quite rare in the data), we expect most classifiers to quickly benefit from increased sample size, and thus \widehat{L}_{CV} to be more conservative.

3 Cross-Validation With Hyper-Parameter Optimization (CVM)

A typical data analysis involves several steps (representing the data, imputation, discretization, variable selection or dimensionality reduction, learning a classifier) each with hundreds of available choices of algorithms in the literature. In addition, each algorithm takes several hyper-parameter values that should be tuned by the user. A general method for tuning the hyper-parameters is to try a set of *predefined* combinations of methods and corresponding values and select the best. We will represent this set with a set \mathbf{a} containing hyper-parameter values, e.g., $\mathbf{a} = \{ \langle \text{no variable selection, K-NN, K=5} \rangle, \langle \text{Lasso, } \lambda = 2, \text{ linear SVM, C=10} \rangle \}$ when the intent is to try K-NN with no variable selection, and a linear SVM using the Lasso algorithm for variable selection. The pseudo-code is shown in Algorithm 2. The symbol $f(x, D, \mathbf{a})$ now denotes the output of the model learned when using hyper-parameters \mathbf{a} on dataset D and applied on input x . Correspondingly, the symbol $f(\cdot, D, \mathbf{a})$ denotes the model produced by applying hyper-parameters \mathbf{a} on D . The quantity $L_{CV}(\mathbf{a})$ is now parameterized by the specific values \mathbf{a} and the minimizer of the loss (maximizer of performance) \mathbf{a}^* is found. The final model returned is $f(\cdot, D, \mathbf{a}^*)$, i.e., the models produced by values \mathbf{a}^* on *all data* D . On one hand, we expect CV with model selection (hereafter, *CVM*) to underestimate performance because estimations are computed using models trained on only a subset of the dataset. On the other hand, we expect *CVM* to overestimate performance because it returns the maximum performance found after trying several hyper-parameter values. In Section 7 we examine this behavior empirically and determine (in concordance with [2],[4]) that indeed when sample size is relatively small

Algorithm 2: K-Fold Cross-Validation with Hyper-Parameter Optimization (Model Selection) $\underline{CVM}(D, \mathbf{a})$
Input: A dataset $D = \{\langle x_i, y_i \rangle, i=1, \dots, N\}$
A set of hyper-parameter value combinations \mathbf{a}
Output: A model M
An estimation of performance (loss) of M

Partition D to K folds F_i

Estimate $\widehat{L_{CV}}(\mathbf{a})$ for each $\mathbf{a} \in \mathbf{a}$:

$$\widehat{e_i}(\mathbf{a}) = \frac{1}{N_i} \sum_{j \in F_i} L(y_j, f(x_j, D_{\setminus i}, \mathbf{a})), \quad \widehat{L_{CV}}(\mathbf{a}) = \frac{1}{K} \sum_{i=1}^K \widehat{e_i}(\mathbf{a})$$

Find minimizer \mathbf{a}^* of $\widehat{L_{CV}}(\mathbf{a})$ // "best hyper-parameters"
 $M = f(\cdot, D, \mathbf{a}^*)$ // the model from all data D with the best hyper-parameters

$\widehat{L_{CVM}} = \widehat{L_{CV}}(\mathbf{a}^*)$
Return $\langle M, \widehat{L_{CVM}} \rangle$

Algorithm 3: $\underline{TT}(D, \mathbf{a})$

Input: A dataset $D = \{\langle x_i, y_i \rangle, i=1, \dots, N\}$
A set of hyper-parameter value combinations \mathbf{a}
Output: A model M
An estimation of performance (loss) of M

Partition D to K folds F_i

Estimate $\widehat{L_{CV}}(\mathbf{a})$ for each $\mathbf{a} \in \mathbf{a}$:

$$\widehat{e_i}(\mathbf{a}) = \frac{1}{N_i} \sum_{j \in F_i} L(y_j, f(x_j, D_{\setminus i}, \mathbf{a})), \quad \widehat{L_{CV}}(\mathbf{a}) = \frac{1}{K} \sum_{i=1}^K \widehat{e_i}(\mathbf{a})$$

Find minimizer \mathbf{a}^* of $\widehat{L_{CV}}(\mathbf{a})$ // "best hyper-parameters"
Find minimizers \mathbf{a}_k of $\widehat{e_k}(\mathbf{a})$ // the minimizers for each fold

Estimate $\widehat{Bias} = \sum_{k=1}^K (\widehat{e_k}(\mathbf{a}^*) - \widehat{e_k}(\mathbf{a}_k))$

$M = f(\cdot, D, \mathbf{a}^*)$, i.e., the model learned on all data D with the best hyper-parameters

$\widehat{L_{TT}} = \widehat{L_{CV}}(\mathbf{a}^*) + \widehat{Bias}$

Return $\langle M, \widehat{L_{TT}} \rangle$

and the number of models in the hundreds CVM overestimates performance. Thus, in these cases other types of estimation protocols are required.

4 The Tibshirani and Tibshirani (TT) Method

The TT method [5] attempts to *heuristically and approximately* estimate the bias of the CV error estimation due to model selection and add it to the final estimate. For

each fold, the bias due to model selection is estimated as $e_k(a^*) - e_k(a_k)$ where, as before, e_k is the average loss in fold k , a_k is the hyper-parameter values that minimizes the loss for fold k , and a^* the global minimizer over all folds. Notice that, if in all folds the same values a_k provide the best performance, then these will also be selected globally and hence $a_k = a^*$ for $k=1, \dots, K$. In this case, *the bias estimate will be zero*. The justification of this estimate for the bias is in [5]. *Notice that CVM and TT return the same model (assuming data are partitioned into the same folds), the one returned by f on all data D using the minimizer a^* of the CV error; however, the two methods return a different estimate of the performance of this model. It is also quite important to notice that TT does not require any additional model training and has minimum computational overhead.*

5 The Nested Cross-Validation Method (NCV)

We could not trace who introduced or coined up first the name Nested Cross-Validation (NCV) method but the authors have independently discovered it and using it since 2005 [6],[13],[14]; one early comment hinting of the method is in [15]. A similar method in a bioinformatics analysis was used as early as 2003 [16]. The main idea is to consider the model selection as part of the learning procedure f . Thus, f tests several hyper-parameter values, selects the best using CV, and returns a *single model*. NCV *cross-validates* f to estimate the performance of the average model returned by f just as normal CV would do with any other learning method taking no hyper-parameters; it's just that f now contains an internal CV trying to select the best model. NCV is a generalization of the Train-Validation-Test protocol where one trains on the Train set for all hyper-parameter values, selects the ones that provide the best performance on Validation, trains on Train+Validation a *single model* using the best-found values and estimates its performance on Test. Since Test is used only once by a single model, performance estimation has no bias due to the model selection process. The final model is trained on *all data* using the best found values for a . NCV generalizes the above protocol to cross-validate every step of this procedure: for each Test, all folds serve as Validation, and this process is repeated for each fold serving as Test. The pseudo-code is shown in Algorithm 4. The pseudo-code is similar to CV (Algorithm 1) with CVM (Cross-Validation with Model Selection, Algorithm 2) serving as the learning function f . *Notice, that NCV returns the same final, single model as CV and TT (assuming the same partition of the data to folds).* Again, the difference regards only the estimation of the performance of this model. NCV requires a quadratic number of models to be trained to the number of folds K (one model is trained for every possible pair of two folds serving as test and validation respectively) and thus it is the most computationally expensive protocol out of the three.

6 Stratification of Folds

In CV folds are partitioned randomly which should maintain *on average* the same class distribution in each fold. However, in small sample size sizes or highly imbal-

Algorithm 4: K-Fold Nested Cross-Validation $\underline{NCV}(D, \mathbf{a})$

Input: A dataset $D = \{\langle x_i, y_i \rangle, i=1, \dots, N\}$

A set of hyper-parameter value combinations \mathbf{a}

Output: A model M

An estimation of performance (loss) of M

Partition D to K folds F_i

$\langle M, \sim \rangle = CVM(D, \mathbf{a})$

Estimation \widehat{L}_{NCV} :

$$\widehat{\epsilon}_i = \frac{1}{N_i} \sum_{j \in F_i} L(y_j, CVM(x_j, D_{\setminus i})) , \quad \widehat{L}_{CV} = \frac{1}{K} \sum_{i=1}^K \widehat{\epsilon}_i$$

Return $\langle M, \widehat{L}_{NCV} \rangle$

anced class distributions it may happen that some folds contain no samples from one of the classes (or in general, the class distribution is very different from the original). In that case, the estimation of performance for that fold will exclude that class. To avoid this case, “in stratified cross-validation, the folds are stratified so that they contain approximately the same proportions of labels as the original dataset” [4]. Notice that leave-one-out CV guarantees that each fold will be unstratified since it contains only one sample which can cause serious estimation problems ([12], p. 151, [4]).

7 Empirical Comparison of Different Protocols

We performed an empirical comparison in order to assess the characteristics of each data-analysis protocol. Particularly, we focus on three specific aspects of the protocol performances: bias, variance of the estimation and the effect of stratification.

7.1 The Experimental Set-Up

Original Datasets: Five datasets from different scientific fields were employed for the experimentations. The computational task for each dataset consists in predicting a binary outcome on the basis of a set of numerical predictors (binary classification). In more detail the **SPECT** [17] dataset contains data from Single Photon Emission Computed Tomography images collected in both healthy and cardiac patients. Data in **Gamma** [18] consist of simulated registrations of high energy gamma particles in an atmospheric Cherenkov telescope, where each gamma particle can be originated from the upper atmosphere (background noise) or being a primary gamma particle (signal). Discriminating biodegradable vs. non-biodegradable molecules on the basis of their chemical characteristics is the aim of the **Biodeg** [19] dataset. The **Bank** [20] dataset was gathered by direct marketing campaigns (phone calls) of a Portuguese banking institution for discriminating customers who want to subscribe a term deposit and those who don’t. Last, **CD4vsCD8** [21] contains the phosphorylation levels of 18 intra-cellular proteins as predictors to discriminate naïve CD4+ and CD8+ human immune system cells. **Table 1** summarizes datasets’ characteristics.

Table 1. Datasets’ characteristics. D_{pool} is a 30% partition from which sub-sampled datasets are produced. $D_{hold-out}$ is the remaining 70% of samples from which an estimation of the true performance is computed.

Dataset Name	# Samples	# Attributes	Classes ratio	$ D_{pool} $	$ D_{hold-out} $	Ref.
SPECT	267	22	3.85	81	186	[17]
Biodeg	1055	41	1.96	317	738	[19]
Gamma	19020	11	1.84	5706	13314	[18]
CD4vsCD8	24126	18	1.13	7238	16888	[21]
Bank	45211	17	7.54	13564	31647	[20]

Sub-Datasets and Hold-out Datasets. Each original dataset D is partitioned into two separate, stratified parts: D_{pool} , containing 30% of the total samples, and the hold-out set $D_{hold-out}$, consisting of the remaining samples. Subsequently, for each D_{pool} 50 sub-datasets are randomly sampled with replacement for each sample size in the set $\{20, 40, 60, 80, 100, 500 \text{ and } 1500\}$, for a total of $5 \times 7 \times 50$ sub-datasets $D_{i,j,k}$ (where i indexes the original dataset, j the sample size, and k the sub-sampling). For sample sizes less than 100 we enforce an equal percentage among the two classes, in order to avoid problems of imbalanced data. Most of the original datasets have been selected with a relatively large sample size so that *each $D_{hold-out}$ is large enough to allow a very accurate (low variance) estimation of performance*. In addition, the size of D_{pool} is also relatively large so that *each sub-sampled dataset to be approximately considered a dataset independently sampled from the data population of the problem*. Nevertheless, we also include a couple of datasets with smaller sample size.

Bias and Variance of each Protocol: For each of the data analysis protocols CVM , TT , and NCV both the stratified and the non-stratified versions are applied to each sub-dataset, in order to select the “best model/hyper-parameter values” and estimate its performance \hat{L} . For each sub-dataset, the same split in $K = 10$ folds was employed for the stratified versions of CVM , TT and NCV , so that the three data-analysis protocols always select exactly the same model, and differ only in the estimation of performance. For the NCV , the internal CV loop uses $K=9$. The bias is computed as $L_{hold-out} - \hat{L}$. Thus, *a positive bias indicates a higher “true” error (i.e., as estimated on the hold-out set) than the one estimated by the corresponding analysis protocol and implies the estimation protocol is optimistic*. For each protocol, original dataset, and sample size the mean bias, its variance and its standard deviation over the 50 sub-samplings are computed and reported in the results’ Tables and Figures below.




Fig. 1. Average loss bias for estimation protocols stratified CVM, TT, and NCV that include model selection. CVM is clearly optimistic and systematically overestimates performance for sample sizes less or equal to 100. TT and NCV do not substantially and systematically overestimate on average, although results vary with dataset.




Fig. 2. Standard deviation of bias over the 50 sub-samplings. CVM has the smallest variance but it overestimates performance. TT and NCV exhibit similar stds.

Performance Metric: All algorithms are presented using a loss function L computed for each sample and averaged out for each fold and then over all folds. The zero-one loss function is typically assumed corresponding to 1-accuracy of the classifier. However, we prefer to use the Area Under the Receiver’s Operating Characteristic Curve (AUC) [22] as the metric of choice for binary classification problems. One advantage is that the AUC does not depend on the prior class distribution. This is necessary in order to pool together estimations stemming from different datasets. In contrast, the zero-one loss depends on the class distribution: for a problem with class distribution of 50-50%, a classifier with accuracy 85% has greatly improved over the baseline of a trivial classifier predicting the majority class; for a problem of 84-16% class distribution, a classifier with 85% accuracy has not improved much over the baseline. Unfortunately, the AUC cannot be expressed as a loss function $L(y, y')$ where y' is a single prediction. Nevertheless, all Algorithms 1-4 remain the same if we substitute $\hat{e}_i = 1 - AUC(f(\cdot, D_{\setminus i}), F_i)$, i.e., the error in fold i is 1 minus the AUC of the model learned by f on all data except fold F_i , as estimated on F_i as the test set.

Model Selection: For generating the hyper-parameter vectors in \mathbf{a} we employed three different modelers: the Logistic Regression classifier ([9], p. 205), as implemented in

Matlab 2013b, that takes no hyper-parameters; the Decision Tree [23], as implemented also in Matlab 2013b with hyper-parameters MinLeaf and MinParents both within $\{1, 2, \dots, 10, 20, 30, 40, 50\}$; Support Vector Machines as implemented in the libsvm software [24] with linear, Gaussian ($\gamma \in \{0.01, 0.1, 1, 10, 100\}$) and polynomial (degree $d \in \{2, 3, 4\}$, $\gamma \in \{0.01, 0.1, 1, 10, 100\}$) kernels, and cost parameter $C \in \{0.01, 0.1, 1, 10, 100\}$. When a classifier takes multiple hyper-parameters, all combinations of choices are tried. Overall, 271 hyper-parameter value combinations and corresponding models are produced each time to select the best.

7.2 Experimental Results

Fig. 1 shows the average loss bias of CVM, TT, and NCV showing that indeed *CVM overestimates performance* for small sample sizes (underestimates error) corroborating the results in [2],[4]. *TT and NCV do not substantially and systematically overestimate, although results vary with dataset.* **Table 2** shows the bias averaged over all datasets, where it is shown that CVM often overestimates the AUC by more than 5 points for small sample sizes. TT seems quite robust with the bias being confined to less than plus or minus 1,7 AUC points for sample sizes more than 20. We perform a t-test for the null hypothesis that the bias is zero, which is typically rejected: all methods usually exhibit some bias whether positive or negative. Nevertheless, in our opinion the bias for TT and NCV is in general acceptable. The non-stratified versions of the protocols have more bias in general. Fig. 2 shows the standard deviation (std) of CVM, TT, and NCV. CVM exhibits the smallest std (and thus variance) but in our opinion it should be avoided since it overestimates performance. **Table 3** shows the averaged std of the bias over all datasets. We apply the O'Brien's modification of Levene's statistical test [25] with the null hypothesis that the variance of a method is the same as the corresponding variance for the same sample size as the NCV. NCV and TT show almost statistically indistinguishable variances. Thus, *within the scope of our experiments and based on the combined analysis of average bias, average variance, and computational complexity the TT protocol seems to be the method of choice.* The non-stratified versions also exhibit slightly larger variance and again, *stratification seems to have only beneficial effects.*

8 Related Work, Discussion and Conclusions

Estimating performance of the final reported model while simultaneously selecting the best pipeline of algorithms and turning their hyper-parameters is a fundamental task for any data analyst. Yet, arguably these issues have not been examined in full depth in the literature. The origins of cross-validation in statistics can be traced back to the “jackknife” technique of Quenouille [26] in the statistical community.

In machine learning, [4] studied the cross-validation without model selection (the title of the paper may be confusing) comparing it against the bootstrap and reaching the important conclusion that (a) CV is preferable to the bootstrap, (b) a value of $K=10$ is preferable for the number of folds versus a leave-one-out, and (c) stratification is also

Table 2. Average Bias over Datasets. P-values produced by a t-test with null hypothesis the mean bias is zero ($P < 0,05^*$, $P < 0,01^{**}$). NS stands for Non-Stratified. CVM systematically overestimates performance. TT and NCV slightly underestimate performance for larger sample sizes.

	CVM	NS-CVM	TT	NS-TT	NCV	NS-NCV
20	0,1551**	0,1892**	0,0525**	0,1142**	-0,0007	-0,032*
40	0,0891**	0,0993**	0,0085	0,0172*	-0,054**	-0,1102**
60	0,0749**	0,0825**	0,0045	0,0136*	0,0083	0,0111
80	0,0507**	0,0563**	-0,0176**	-0,0097	-0,0228**	-0,0338**
100	0,0681**	0,0731**	-0,0036	0,0079	0,0131*	0,0183**
500	0,0072**	0,0073**	-0,025**	-0,0261**	-0,0054*	-0,0055*
1500	-0,0005	0,0002	-0,0139**	-0,0136**	-0,0034**	-0,003*

Table 3. Average bias STDs over Datasets. P-values produced by a test with null hypothesis that the variances are the same as the corresponding variance of the NCV protocol ($P < 0,05^*$, $P < 0,01^{**}$). NS stands for Non-Stratified. NCV and TT have indistinguishable variances and are conservative. CVM has smaller variance but overestimates performance.

	CVM	NS-CVM	TT	NS-TT	NCV	NS-NCV
20	0,1134**	0,119**	0,1826	0,1989**	0,1616	0,2055**
40	0,0751**	0,0808**	0,1194**	0,1298*	0,1571	0,1784
60	0,0659*	0,0727	0,0916	0,1017	0,0828	0,0936
80	0,0497**	0,0526**	0,0757	0,0804	0,0731	0,0871*
100	0,0651*	0,0697	0,0917	0,0903	0,0826	0,0845
500	0,028	0,0237	0,0341	0,0351	0,0285	0,0308
1500	0,0119	0,014	0,018**	0,0192**	0,0125	0,0143

always preferable. In terms of theory, Bengio [11] showed that there exist no unbiased estimator for the variance of the CV performance estimation, which impact hypothesis testing of performance using the CV.

To the extent of our knowledge the first to study the problem of bias *in the context of model selection* in machine learning is [3]. Varma [27] demonstrated the optimism of the CVM protocol and instead suggests the use of the NCV protocol. Unfortunately, all their experiments are performed on simulated data only. Tibshirani and Tibshirani [5] introduced the TT protocol but unfortunately they do not compare it against alternatives and they include only a proof-of-concept experiment on a single dataset. Thus, the present paper is the first work that compares all three protocols (CVM, NCV, and TT) on multiple real datasets.

Based on our experiments we found evidence that the TT method is unbiased (slightly conservative) for sample sizes above 20, has about the same variance as the NCV (the other conservative alternative), and does not introduce additional computational overhead. Within the scope of our experiments, *we would thus suggest analysts to employ the TT estimation protocol*. In addition, we corroborate the results in [4] that stratification exhibits smaller estimation variance and we encourage its use.

We would also like to acknowledge the limited scope of our experiments in terms of the number and type of datasets, the inclusion of other steps into the analysis (such as variable selection), the inclusion of other procedures for hyper-parameter optimization that dynamically decide to consider value combinations, varying the number of folds K in the protocols, using other performance metrics, experimentation with regression methods and others which form our future work on the subject in order to obtain definite answers to these research questions.

We also note *the concerning issue* that the variance of estimation for small sample sizes is large, again in concordance with the experiments in [2]. The authors in the latter advocate methods that may be biased but exhibit reduced variance. However, we believe that CVM is too biased no matter its variance; implicitly the authors in [2] agree when they declare that model selection should be integrated in the performance estimation procedure in such a way that test samples are never employed for selecting the best model. Instead, they suggest as alternatives limiting the extent of the search of the hyper-parameters or performing model averaging. In our opinion, neither option is satisfactory for all analysis purposes and more research is required.

9 Acknowledgements

This work was conducted partially in the framework of the BIOSYS research project, Action KRIPIS, project No MIS-448301 (2013SE01380036) by the Greek General Secretariat for Research and Technology. It was funded partially by the STATegra EU FP7 project, No 306000.

10 References

1. Anguita, D., Ghio, A., Oneto, L., Ridella, S.: In-Sample and Out-of-Sample Model Selection and Error Estimation for Support Vector Machines. *IEEE Trans. Neural Networks Learn. Syst.* 23, 1390–1406 (2012).
2. Cawley, G.C., Talbot, N.L.C.: On Over-fitting in Model Selection and Subsequent Selection Bias in Performance Evaluation. *J. Mach. Learn. Res.* 11, 2079–2107 (2010).
3. Jensen, D.D., Cohen, P.R.: Multiple comparisons in induction algorithms. *Mach. Learn.* 38, 309–338 (2000).
4. Kohavi, R.: A Study of Cross-Validation and Bootstrap for Accuracy Estimation and Model Selection. *International Joint Conference on Artificial Intelligence*. pp. 1137–1143 (1995).
5. Tibshirani, R.J., Tibshirani, R.: A bias correction for the minimum error rate in cross-validation. *Ann. Appl. Stat.* 3, 822–829 (2009).
6. Statnikov, A., Aliferis, C.F., Tsamardinos, I., Hardin, D., Levy, S.: A comprehensive evaluation of multicategory classification methods for microarray gene expression cancer diagnosis. *Bioinformatics*. 21, 631–43 (2005).
7. Duda, R.O., Hart, P.E., Stork, D.G.: *Pattern Classification* (2nd Edition). (2000).
8. Mitchell, T.M.: *Machine Learning*. (1997).
9. Bishop, C.M.: *Pattern Recognition and Machine Learning* (Information Science and Statistics). (2006).

10. Hastie, T., Tibshirani, R., Friedman, J.: The Elements of Statistical Learning. Elements. 1, 337–387 (2009).
11. Bengio, Y., Grandvalet, Y.: Bias in Estimating the Variance of K-Fold Cross-Validation. Statistical Modeling and Analysis for Complex Data Problem. pp. 75–95 (2005).
12. Witten, I.H., Frank, E.: Data Mining: Practical Machine Learning Tools and Techniques, Second Edition (Morgan Kaufmann Series in Data Management Systems). (2005).
13. Lagani, V., Tsamardinos, I.: Structure-based variable selection for survival data. Bioinformatics. 26, 1887–1894 (2010).
14. Statnikov, A., Tsamardinos, I., Dosbayev, Y., Aliferis, C.F.: GEMS: a system for automated cancer diagnosis and biomarker discovery from microarray gene expression data. Int. J. Med. Inform. 74, 491–503 (2005).
15. Salzberg, S.: On Comparing Classifiers : Pitfalls to Avoid and a Recommended Approach. Data Min. Knowl. Discov. 328, 317–328 (1997).
16. Iizuka, N., Oka, M., Yamada-Okabe, H., Nishida, M., Maeda, Y., Mori, N., Takao, T., Tamesa, T., Tangoku, A., Tabuchi, H., Hamada, K., Nakayama, H., Ishitsuka, H., Miyamoto, T., Hirabayashi, A., Uchimura, S., Hamamoto, Y.: Oligonucleotide microarray for prediction of early intrahepatic recurrence of hepatocellular carcinoma after curative resection. Lancet. 361, 923–9 (2003).
17. Kurgan, L.A., Cios, K.J., Tadeusiewicz, R., Ogiela, M., Goodenday, L.S.: Knowledge discovery approach to automated cardiac SPECT diagnosis. Artif. Intell. Med. 23, 149–69 (2001).
18. Bock, R.K., Chilingarian, A., Gaug, M., Hakl, F., Hengstebeck, T., Jiřina, M., Klaschka, J., Kotrč, E., Savický, P., Towers, S., Vaiciulis, A., Wittek, W.: Methods for multidimensional event classification: a case study using images from a Cherenkov gamma-ray telescope. Nucl. Instruments Methods Phys. Res. Sect. A Accel. Spectrometers, Detect. Assoc. Equip. 516, 511–528 (2004).
19. Mansouri, K., Ringsted, T., Ballabio, D., Todeschini, R., Consonni, V.: Quantitative structure-activity relationship models for ready biodegradability of chemicals. J. Chem. Inf. Model. 53, 867–78 (2013).
20. Moro, S., Laureano, R.M.S.: Using Data Mining for Bank Direct Marketing: An application of the CRISP-DM methodology. Eur. Simul. Model. Conf. 117–121 (2011).
21. Bendall, S.C., Simonds, E.F., Qiu, P., Amir, E.D., Krutzik, P.O., Finck, R., Bruggner, R. V., Melamed, R., Trejo, A., Ornatsky, O.I., Balderas, R.S., Plevritis, S.K., Sachs, K., Pe'er, D., Tanner, S.D., Nolan, G.P.: Single-cell mass cytometry of differential immune and drug responses across a human hematopoietic continuum. Science. 332, 687–96 (2011).
22. Fawcett, T.: An introduction to ROC analysis, (2006).
23. Coppersmith, D., Hong, S.J., Hosking, J.R.M.: Partitioning Nominal Attributes in Decision Trees. Data Min. Knowl. Discov. 3, 197–217 (1999).
24. Chang, C.-C., Lin, C.-J.: LIBSVM: a library for support vector machines. ACM Trans. Intell. Syst. 2, 1–39 (2011).
25. O'brien, R.G.: A General ANOVA Method for Robust Tests of Additive Models for Variances. J. Am. Stat. Assoc. 74, 877–880 (1979).
26. Quenouille, M.H.: Approximate tests of correlation in time-series 3, (1949).
27. Varma, S., Simon, R.: Bias in error estimation when using cross-validation for model selection. BMC Bioinformatics. 7, 91 (2006).